



The psql Client			Bulk Loads & Extracts with /copy	
psql is PostgreSQL's Command Line (non-GUI) client utility, it operates in both interactive (query) and non-interactive (batch) modes.			The /copy meta command is the psql utility's mechanism to bulk load and/or create extracts of database tables.	
Common Command Line Arguments	Common Meta Commands	Common Describe Meta Commands	<pre>\COPY table_name [ ( column_name [ , ... ] ) ] FROM { 'filename'         PROGRAM 'command'         STDIN } [ [ WITH ] ( option [ , ... ] ) ]  \COPY { table_name [ ( column_name [ , ... ] ) ]         ( query ) } TO { 'filename'        PROGRAM 'command'        STDOUT } [ [ WITH ] ( option [ , ... ] ) ]  where option can be one of:  FORMAT format_name OIDS [ boolean ] FREEZE [ boolean ] DELIMITER 'delimiter_character' NULL 'null_string' HEADER [ boolean ] QUOTE 'quote_character' ESCAPE 'escape_character' FORCE_QUOTE { ( column_name [ , ... ] ) ! * } FORCE_NOT_NULL ( column_name [ , ... ] ) FORCE_NULL ( column_name [ , ... ] ) ENCODING 'encoding_name'</pre>	
<b>General options</b> <b>-c, --command=COMMAND</b> Run only single command (SQL or internal) and exit <b>-d, --dbname=DBNAME</b> Database name to connect to (default: "student") <b>-f, --file=FILENAME</b> Execute commands from file, then exit <b>-l, --list</b> List available databases, then exit <b>-V, --version</b> output version information, then exit <b>-X, --no-psqlrc</b> do not read startup file (~/.psqlrc) <b>-1 ("one"), --single-transaction</b> execute as a single transaction (if non-interactive) <b>-?, --help</b> show this help, then exit <b>Input and output options</b> <b>-a, --echo-all</b> echo all input from script <b>-e, --echo-queries</b> echo commands sent to server <b>-E, --echo-hidden</b> display queries that internal commands generate <b>-L, --log-file=FILENAME</b> send session log to file <b>-o, --output=FILENAME</b> send query results to file (or  pipe) <b>-s, --single-step</b> single-step mode (confirm each query) <b>Output format options</b> <b>-A, --no-align</b> unaligned table output mode <b>-F, --field-separator=STRING</b> field separator for unaligned output <b>-H, --html</b> HTML table output mode <b>-P, --pset=VAR[=ARG]</b> set printing option VAR to ARG (see \pset command) <b>-R, --record-separator=STRING</b> record separator for unaligned output <b>-t, --tuples-only</b> print rows only <b>-T, --table-attr=TEXT</b> set HTML table tag attributes (e.g., width, border) <b>-x, --expanded</b> turn on expanded table output <b>-z, --field-separator-zero</b> set field separator for unaligned output to zero byte <b>-0, --record-separator-zero</b> set record separator for unaligned output to zero byte <b>Connection options</b> <b>-h, --host=HOSTNAME</b> database server host or socket directory (PGPASSWORD) <b>-p, --port=PORT</b> database server port (PGPORT) <b>-U, --username=USERNAME</b> database user name (PGUSER) <b>-W, --password</b> force password prompt (PGPASSWORD)	<b>General</b> <b>\q [FILE] or :</b> execute query (and send results to file or  pipe) <b>\h [NAME]</b> help on syntax of SQL commands, * for all commands <b>\q</b> quit psql <b>\watch [SEC]</b> execute query every SEC seconds <b>Query Buffer</b> <b>\e [FILE] [LINE]</b> edit the query buffer (or file) with external editor <b>\ef [FUNCNAME [LINE]]</b> edit function definition with external editor <b>\p</b> show the contents of the query buffer <b>\r</b> reset (clear) the query buffer <b>\s [FILE]</b> display history or save it to file <b>\w FILE</b> write query buffer to file <b>Input/Output</b> <b>\i FILE</b> execute commands from file <b>\ir FILE</b> as \i, but relative to location of current script <b>\o [FILE]</b> send all query results to file or  pipe <b>Formatting</b> <b>\a</b> toggle between unaligned and aligned output mode <b>\f [STRING]</b> show or set field separator for unaligned query output <b>\t [on off]</b> show only rows <b>\x [on off auto]</b> toggle expanded output <b>Connection</b> <b>\c[connect] [(DBNAME - USER - HOST - PORT)-]   conninfo)</b> connect to new database <b>\encoding [ENCODING]</b> show or set client encoding <b>\password [USERNAME]</b> securely change the password for a user <b>\conninfo</b> display information about current connection <b>Operating System</b> <b>\cd [DIR]</b> change the current working directory <b>\setenv NAME [VALUE]</b> set or unset environment variable <b>\timing [on off]</b> toggle timing of commands (currently off) <b>\i [COMMAND]</b> execute command in shell or start interactive shell	<b>S = show system objects</b> <b>+ = additional detail</b> <b>\d[S+]</b> list tables, views, and sequences <b>\d[S+] NAME</b> describe table, view, sequence, or index <b>\sf[+] FUNCNAME</b> show a function's definition <b>The following accept pattern/wildcards</b> <b>\da[S]</b> list aggregates <b>\db[+]</b> list tablespaces <b>\ddp</b> list default privileges <b>\det[+]</b> list foreign tables <b>\des[+]</b> list foreign servers <b>\deu[+]</b> list user mappings <b>\dew[+]</b> list foreign-data wrappers <b>\dff[S+]</b> list functions <b>\dfa[S+]</b> list only aggregate functions <b>\dfn[S+]</b> list only normal functions <b>\dft[S+]</b> list only trigger functions <b>\dfw[S+]</b> list only window functions <b>\dgr[+]</b> list roles <b>\di[S+]</b> list indexes <b>\dl[S+]</b> list procedural languages <b>\dm[S+]</b> list materialized views <b>\dn[S+]</b> list schemas <b>\do[S]</b> list operators <b>\dO[S+]</b> list collations <b>\dp</b> list table, view, and sequence access privileges <b>\drds</b> list per-database role settings <b>\ds[S+]</b> list sequences <b>\dt[S+]</b> list tables <b>\dT[S+]</b> list data types <b>\du[+]</b> list roles <b>\dv[S+]</b> list views <b>\dE[S+]</b> list foreign tables <b>\dx[+]</b> list extensions <b>\dy</b> list event triggers <b>\d[+]</b> list databases		
<b>Using Conninfo Strings (Common Values)</b> Conninfo strings are used to specify connection parameters using a single string. Example: psql "host=pghost port=5432 dbname=payroll user=student password=pass connect_timeout=10"				<b>The pg_settings Table</b> The pg_setting table provides information about current database settings as well as their descriptions.
<b>host</b> Name of host to connect to. <b>hostaddr</b> Numeric IP address of host to connect to <b>port</b> Port number to connect to <b>dbname</b> The database name <b>user</b> User name <b>password</b> Password <b>application_name</b> Application name (shows in pg_stat_activity)	<b>client_encoding</b> Client character encoding <b>connect_timeout</b> Connection timeout (seconds) <b>sslmode</b> One of six values (disable, allow, prefer (default), require, verify-ca, verify-full) <b>sslcert</b> Path to SSL Client Certificate <b>sslkey</b> Path to SSL Key file <b>sslrootcert</b> Path to SSL Root certificate			



The pg_dump Client		PostgreSQL Training & Consulting
pg_dump dumps a database as a text file or to other formats. This command shares connection options with the pg_dump command.		<b>Open Technology Group, Inc.</b> <a href="http://www.otg-nc.com">http://www.otg-nc.com</a> <a href="mailto:info@otg-nc.com">info@otg-nc.com</a> <b>877-258-8987</b>
Common Command Line Arguments	The pg_restore Client	Statistics Tables (cont'd)
<p><b>General options</b></p> <p>-f, --file=FILENAME output file or directory name</p> <p>-F, --format=cldt t output file format (custom, directory, tar, plain text (default))</p> <p>-j, --jobs=NUM use this many parallel jobs to dump</p> <p>-v, --verbose verbose mode</p> <p>-V, --version output version information, then exit</p> <p>-Z, --compress=0-9 compression level for compressed formats</p> <p>--lock-wait-timeout=TIMEOUT fail after waiting TIMEOUT for a table lock</p> <p>?, --help show this help, then exit</p> <p><b>Options controlling the output content</b></p> <p>-a, --data-only dump only the data, not the schema</p> <p>-b, --blobs include large objects in dump</p> <p>-c, --clean clean (drop) database objects before recreating</p> <p>-C, --create include commands to create database in dump</p> <p>-E, --encoding=ENCODING dump the data in encoding ENCODING</p> <p>-n, --schema=SCHEMA dump the named schema(s) only</p> <p>-N, --exclude-schema=SCHEMA do NOT dump the named schema(s)</p> <p>-o, --oids include OIDs in dump</p> <p>-O, --no-owner skip restoration of object ownership in plain-text format</p> <p>-s, --schema-only dump only the schema, no data</p> <p>-S, --superuser=NAME superuser user name to use in plain-text format</p> <p>-t, --table=TABLE dump the named table(s) only</p> <p>-T, --exclude-table=TABLE do NOT dump the named table(s)</p> <p>-x, --no-privileges do not dump privileges (grant/revoke)</p> <p>--binary-upgrade for use by upgrade utilities only</p> <p>--column-inserts dump data as INSERT commands with column names</p> <p>--disable-dollar-quoting disable dollar quoting, use SQL standard quoting</p> <p>--disable-triggers disable triggers during data-only restore</p> <p>--exclude-table-data=TABLE do NOT dump data for the named table(s)</p> <p>--if-exists use IF EXISTS when dropping objects</p> <p>--inserts dump data as INSERT commands, rather than COPY</p> <p>--no-security-labels do not dump security label assignments</p> <p>--no-synchronized-snapshots do not use synchronized snapshots in parallel jobs</p> <p>--no-tablespaces do not dump tablespace assignments</p> <p>--no-unlogged-table-data do not dump unlogged table data</p> <p>--quote-all-identifiers quote all identifiers, even if not key words</p> <p>--section=SECTION dump named section (pre-data, data, or post-data)</p> <p>--serializable-deferrable wait until the dump can run without anomalies</p> <p>--use-set-session-authorization use SET SESSION AUTHORIZATION commands instead of ALTER OWNER commands to set ownership</p> <p><b>Connection Options</b></p> <p>Connection options match those of the psql command.</p>	<p>The pg_restore client is used to restore backups made with pg_dump using the tar, directory, or custom dump format. It creates a <i>plain text</i> dump from an existing custom dump (suitable for use with psql); or connects directly to PostgreSQL to restore data.</p> <p>Connection arguments match that of psql.</p> <p><b>General options</b></p> <p>-d, --dbname=NAME connect to database name</p> <p>-f, --file=FILENAME output file name</p> <p>-l, --list print summarized TOC of the archive</p> <p>-v, --verbose verbose mode</p> <p>-V, --version output version information, then exit</p> <p>?, --help show this help, then exit</p> <p><b>Options controlling the restore</b></p> <p>Options match the output options of the pg_dump command.</p> <p><b>Connection Options</b></p> <p>Connection options match those of the psql command.</p> <p><b>Custom Dump Formats</b></p> <p>Custom dump formats (-F) allow pg_dump to generate improved format outputs that allow for more flexible recovery, and (in some cases) faster output.</p> <p><b>-Fd Format: directory</b></p> <p>Produces the output in a directory (rather than a single file.) When using -Fd, the -f file argument should be a directory. Can be combined with the -j argument (and usually --serializable-deferrable) to produce a <i>parallel</i> dump to dump multiple objects simultaneously.</p> <p><b>-Fc Format: custom</b></p> <p>Produces a special compressed format. Optimal for single object/table restores with pg_restore.</p> <p><b>-Ft Format: tar</b></p> <p>Produces a tar file output that matches the output that would result if the -Fd argument was used, then the result tarred up. Cannot be used with -j for parallel dumps.</p> <p><b>Statistics Tables</b></p> <p><b>pg_stat_activity</b> Provide information about currently running queries/processes on the server (cluster-wide.)</p> <p><b>pg_stat_bgwriter</b> Provide cluster-wide statistics from the background writer process, including checkpoint statistics and buffer write statistics.</p> <p><b>pg_stat_database</b> Provide rollup statistics at the database level</p> <p><b>pg_stat_database_conflicts</b> Provide database-level statistics regarding the number of queries that have been canceled in this database due to dropped tablespaces, lock timeouts, old snapshots, pinned buffers and deadlocks (standby servers only.)</p> <p><b>pg_stat_replication</b> Provide statistic information regarding replication slave servers and their current state.</p>	<p><b>Statistics Tables (cont'd)</b></p> <p><b>pg_stat_all_tables</b> Provides row-level statistics for each table in the current database.</p> <p><b>pg_stat_sys_tables</b> Same as pg_stat_all_tables, except that only system tables are shown.</p> <p><b>pg_stat_user_tables</b> Same as pg_stat_all_tables, except that only user tables are shown.</p> <p><b>pg_stat_xact_all_tables</b> Similar to pg_stat_all_tables, but counts actions taken so far within the current transaction (which are not yet included in pg_stat_all_tables and related views).</p> <p><b>pg_stat_xact_sys_tables</b> Same as pg_stat_xact_all_tables, except that only system tables are shown.</p> <p><b>pg_stat_xact_user_tables</b> Same as pg_stat_xact_all_tables, except that only user tables are shown.</p> <p><b>pg_stat_all_indexes</b> Provides row-level statistics for each index in the current database.</p> <p><b>pg_stat_sys_indexes</b> Same as pg_stat_all_indexes, except that only indexes on system tables are shown.</p> <p><b>pg_stat_user_indexes</b> Same as pg_stat_all_indexes, except that only indexes on user tables are shown.</p> <p><b>pg_statio_all_tables</b> Provides block-level statistics for each table in the current database.</p> <p><b>pg_statio_sys_tables</b> Same as pg_statio_all_tables, except that only system tables are shown.</p> <p><b>pg_statio_user_tables</b> Same as pg_statio_all_tables, except that only user tables are shown.</p> <p><b>pg_statio_all_indexes</b> Provides block-level statistics for each index in the current database.</p> <p><b>pg_statio_sys_indexes</b> Same as pg_statio_all_indexes, except that only indexes on system tables are shown.</p> <p><b>pg_statio_user_indexes</b> Same as pg_statio_all_indexes, except that only indexes on user tables are shown.</p> <p><b>pg_statio_all_sequences</b> Provides block-level statistics for each sequence in the current database.</p> <p><b>pg_statio_sys_sequences</b> Same as pg_statio_all_sequences, except that only system sequences are shown.</p> <p><b>pg_statio_user_sequences</b> Same as pg_statio_all_sequences, except that only user sequences are shown.</p> <p><b>pg_stat_user_functions</b> Provides information about function execution (number of calls, total time, and the amount of time spent in the function itself (self_time))</p> <p><b>pg_stat_xact_user_functions</b> Similar to pg_stat_user_functions, but counts only calls during the current transaction (which are not yet included in pg_stat_user_functions).</p>